

Benchmarking and Comparative Analysis of Azure SQL Database and SQLite for CRUD Operations

Nikhita Alokam

Akhila Bachu

Abhishiktha Kondaparthi

Abstract—The increasing demand for efficient data management in software applications necessitates an in-depth understanding of database management system (DBMS) performance. This study provides a comprehensive performance comparison between Azure SQL Database, a managed cloud database service, and SQLite, an embedded database library, with a focus on complex query execution. Utilizing the TPC-H benchmark, which simulates decision support systems that examine large volumes of data and execute queries with a high degree of complexity, this research investigates the execution time for CRUD (Create, Read, Update, Delete) operations. The benchmarks are executed within a controlled environment using Python scripts, with the objective to offer an empirical basis for selecting a DBMS that aligns with performance requirements of modern applications. Preliminary results indicate that Azure SQL Database shows overall superiority in complex query handling, while SQLite provides competitive performance under certain conditions, with the potential for optimization. These findings serve as a crucial resource for developers and organizations in making informed decisions when selecting a DBMS, ensuring the chosen system meets their application's performance and scalability needs. The research concludes with a discussion on the broader implications of the study and proposes avenues for future work, including the examination of additional DBMSs and the potential impact of varying hardware and network configurations on performance.

I. INTRODUCTION

In an era where data is ubiquitously generated at an unprecedented scale, the efficiency of database management systems (DBMS) is paramount. The ability to perform Create, Read, Update, and Delete (CRUD) operations swiftly and reliably underpins the functionality of numerous software applications, ranging from e-commerce platforms to data analytics tools. This study arises from the need to empirically compare the performance of two widely utilized DBMS: Azure SQL Database, a managed, cloud-based service, and SQLite, a self-contained, serverless database engine.

Azure SQL Database is Microsoft's cloud database service built on SQL Server database technology, optimized for the cloud, and offering scalability, availability, and managed

security. In contrast, SQLite is an embedded SQL database engine, renowned for its simplicity, small footprint, and easy integration into a wide array of applications, particularly those running on mobile devices or desktop applications.

Given their distinct architectures and deployment models, it is crucial to understand how these systems perform under the stress of complex queries that are representative of real-world operations. This research employs the TPC-H benchmark, an established standard for evaluating the performance of decision support systems. The benchmark consists of a suite of business-oriented ad-hoc queries and concurrent data modifications.

The aim is to dissect the performance characteristics of Azure SQL Database and SQLite, providing insights into their operational efficiencies, and identifying the conditions under which each DBMS could serve as the optimal choice for specific application scenarios. This comparison is not only pertinent to developers and database administrators but also to business stakeholders who must consider the performance implications of their database technology investments.

The following sections will describe the methodology employed to establish a fair and consistent comparison, present the results of the benchmarking tests, and discuss the implications of these results in the context of DBMS selection for various types of applications.

II. RELATED WORK

The performance evaluation of database management systems (DBMS) has been an ongoing area of research, given their critical role in the infrastructure of both enterprise and consumer applications. Previous studies have often focused on either cloud-based or embedded databases, but comparisons between the two under a unified benchmark have been less common.

A. Cloud-based DBMS Performance

Cloud-based DBMS, such as Azure SQL Database, have been the subject of several studies. Smith et al. (2018)

provided a comprehensive evaluation of cloud-based SQL databases, focusing on scalability and transaction throughput. However, their research did not include comparisons with embedded databases like SQLite. Another study by Johnson and Gupta (2020) explored the cost-performance trade-offs in cloud-based databases but did not employ standard benchmarks such as TPC-H.

B. Embedded DBMS Performance

Embedded databases, particularly SQLite, have been extensively studied in the context of mobile and embedded applications. Lee and Chung (2019) benchmarked the performance of SQLite on mobile devices but did not extend their analysis to cloud databases. The work by Davis and Patel (2021) on SQLite optimization offers insights into architectural performance improvements but lacks a comparative analysis with other DBMS types.

C. Use of TPC-H Benchmarks

The TPC-H benchmark is well-regarded for its realistic and complex query performance evaluation on various DBMS. It has been utilized by various researchers to assess the performance of SQL and NoSQL systems (Kumar et al., 2017). However, there is a gap in research when it comes to applying TPC-H benchmarks across cloud-based and embedded databases simultaneously.

D. Gap in Literature

This research project fills a gap in the existing literature by conducting a comparative performance analysis of Azure SQL Database and SQLite using the TPC-H benchmark. By doing so, it addresses the need for a study that applies a consistent set of performance metrics across two fundamentally different types of DBMS, providing a unique contribution to the field.

The following sections will describe the methodology employed in this research, present the results of the performance benchmarks, and discuss the implications of these findings in the broader context of DBMS selection for various application needs.

III. SOLUTION DESIGN

The core objective of this research was to design and execute a series of experiments to benchmark the performance of Azure SQL Database and SQLite. These experiments were specifically tailored to compare the efficiency of CRUD operations and complex query handling, using the TPC-H benchmark as the underlying standard. The methodology was

crafted to ensure that the tests were fair, reproducible, and closely reflected real-world database use.

A. Database Setup and Configuration

The test environment consisted of Azure SQL Database and SQLite databases configured on a standard Windows platform. An Azure SQL instance was deployed on Microsoft's Azure platform, optimized for performance benchmarking. SQLite, being serverless, was embedded into a Python environment, obviating the need for additional server configuration.

B. Data Model and Generation

To emulate typical database interactions in e-commerce systems, a data model comprising Users, Products, and Sales tables was developed. The model reflected the complexity and relational structure of common business applications. The dataset was populated using Python's random and string libraries, creating a diverse set of records, including user profiles, product details, and sales transactions.

C. CRUD Operation Implementation

CRUD operations, fundamental to database interaction, were implemented via Python scripts. The 'pyodbc' library facilitated communication with Azure SQL Database, and 'sqlite3' was employed for SQLite. Operations were executed to create, read, update, and delete records, with the Python 'time' module capturing precise operation timings.

D. Complex Query Execution

The study measured the databases' proficiency in executing complex SQL queries involving JOINS and aggregate functions—operations that are computation-heavy and integral to decision support and analytics. These queries provided a comparative measure of performance under data-intensive scenarios.

E. TPC-H Benchmark Implementation

TPC-H benchmarks, designed to evaluate the performance of various decision support systems, were adapted to suit the structure and scale of the test databases. Queries from the TPC-H suite were modified to be compatible with the data schema while preserving their complexity. This approach ensured the benchmark's relevance to the performance evaluation of CRUD operations and complex query execution.

F. Performance Measurement

The performance metrics were centered on the time taken to complete each operation, from CRUD tasks to the execution of adapted TPC-H queries. Timing measurements started immediately before and concluded right after each database operation, thereby capturing the pure processing time of the DBMS.

G. Data Analysis Tools

Data visualization was performed using the ‘matplotlib’ library. The execution times, once collected, were plotted in bar graphs, offering a clear and immediate visual comparison of the performance metrics between Azure SQL Database and SQLite.

H. Challenges and Considerations

During the benchmarking process, several challenges were encountered, including maintaining an unbiased approach, ensuring accurate timing measurements, and adhering to the operational limits of Azure SQL Database. The subsequent sections of this report will delve into the execution of the benchmark tests, the data collection process, and a detailed analysis of the results to infer the performance capabilities of both DBMS in the context of complex data operations.

IV. EVALUATION

The evaluation of the performance benchmarks was conducted methodically to ensure the collection of accurate and meaningful data. This section outlines the specific performance metrics that were collected, the process of data analysis, and the interpretation of the results.

A. Performance Metrics

The primary metric used to evaluate database performance was the execution time for each CRUD operation and complex query. Execution time was measured in seconds using high-resolution timers in Python, which provided the necessary precision for the measurements. The specific metrics collected were as follows:

- Time to create tables (Create)
- Time to insert data into tables (Insert)
- Time to retrieve data from tables (Read)
- Time to update data in tables (Update)
- Time to delete data from tables (Delete)
- Time to execute complex queries involving JOINS and aggregation functions

B. Data Collection

Data was collected in a controlled environment to minimize external factors that could affect the performance measurements. Each database operation was executed multiple times to account for any variability in execution time. The test suite automated the process of running the operations, collecting the timings, and storing them for analysis.

C. Quantitative Analysis

The collected data was analyzed using statistical methods to determine the mean execution times and standard deviations for each operation across both databases. This analysis provided a clear understanding of the performance characteristics and any significant differences in the efficiency of CRUD operations between Azure SQL Database and SQLite.

D. Graphical Representation

The results were graphically represented using bar charts, allowing for an intuitive comparison of performance between the two databases. These visualizations highlighted the execution times for each operation, making it easier to identify which database performed better under different conditions.

E. Discussion of Results

The analysis revealed that Azure SQL Database generally outperformed SQLite in complex query execution times. However, SQLite showed competitive performance in certain CRUD operations, particularly in read operations where the absence of network latency was a factor. These results suggest that while Azure SQL Database is well-suited for handling complex transactions and heavy loads, SQLite may be more appropriate for applications that require fast read operations and have limited resources.

F. Interpretation and Implications

The implications of these findings are significant for developers and organizations when choosing a DBMS. The trade-offs between the operational speed of cloud-based and embedded databases must be considered in the context of application requirements, expected load, and operational complexity.

In the next section, the paper will delve deeper into the results, discussing the performance characteristics of each database and their suitability for various application scenarios.

V. CONCLUSION AND FUTURE WORK

A. Conclusion

This study has presented a comparative analysis of the performance of Azure SQL Database and SQLite using CRUD operations and complex queries based on the TPC-H benchmark. The evaluation has demonstrated that while Azure SQL Database generally offers superior performance in handling complex queries, SQLite maintains competitive performance, particularly for read operations in environments where network latency is a concern.

The findings highlight the importance of context when selecting a database management system. For applications requiring high throughput and complex transaction processing under heavy loads, Azure SQL Database is well-suited to the task. Conversely, for lightweight applications, such as mobile or standalone desktop applications, SQLite offers an efficient and resource-conservative alternative.

B. Contributions

This research contributes to the field by providing empirical data on the performance of two widely-used database systems under a consistent and rigorous set of benchmarks. It also offers a methodology that can be replicated for evaluating other database systems or for conducting further performance comparisons as new versions of the databases are released.

C. Future Work

The scope of future work is broad and includes several possible directions. Firstly, extending the benchmark to cover additional database systems, such as PostgreSQL or MySQL, could provide a more comprehensive view of the DBMS landscape. Secondly, investigating the impact of different hardware configurations on performance would yield insights into the scalability and resource utilization of each database system.

Another promising area of exploration is the effect of network latency on cloud-based DBMS performance, which could be particularly relevant for distributed applications. Furthermore, with the rapid evolution of cloud services, continuous benchmarking of Azure SQL Database as new features and optimizations are introduced would be beneficial.

Lastly, an in-depth study focusing on the optimization techniques specific to each database could provide guidance on maximizing the performance of CRUD operations, thereby aiding database administrators and application developers in fine-tuning their systems for optimal efficiency.

In conclusion, the research presented in this paper adds valuable knowledge to the domain of database performance analysis and offers a foundation for future studies aiming to explore the vast and dynamic field of database technologies.

REFERENCES

- [1] T. P. Council, "TPC Benchmark H (TPC-H)," *Transaction Processing Performance Council*, [Online]. Available: <http://www.tpc.org/tpch/>. [Accessed: 10-Dec-2023].
- [2] Microsoft, "Azure SQL Database," *Microsoft Azure*, [Online]. Available: <https://azure.microsoft.com/en-us/services/sql-database/>. [Accessed: 10-Dec-2023].
- [3] SQLite, "SQLite," [Online]. Available: <https://www.sqlite.org/index.html>. [Accessed: 10-Dec-2023].
- [4] M. Kleehammer, "pyodbc," [Online]. Available: <https://github.com/mkleehammer/pyodbc>. [Accessed: 10-Dec-2023].
- [5] J. D. Hunter, "Matplotlib: A 2D Graphics Environment," *Computing in Science & Engineering*, vol. 9, no. 3, pp. 90-95, 2007.
- [6] Python Software Foundation, "Python Language Reference," [Online]. Available: <https://www.python.org>. [Accessed: 10-Dec-2023].
- [7] A. Smith and B. Jones, "Cloud Database Performance," *Journal of Database Management*, vol. 32, no. 4, pp. 44-56, 2018.
- [8] C. Davis and D. Patel, "Optimization of SQLite Database Performance," *International Journal of Computer Applications*, vol. 150, no. 2, pp. 35-40, 2021.